

Scheme 4

Stage 3 Retrieval and Manual Relevance Triage Scheme
Full-text availability, retrieval need, and adjudication checklist

DRAFT FOR EXPERT EVALUATION. These coding schemes are a working draft circulated for expert evaluation. They have not been applied to a final review corpus. The current manuscript is a test run built with an earlier, coarser generation of these schemes; the refinements proposed here are awaiting expert sign-off before any re-run.

At a Glance

Workflow position

Bridge between Stage 2 coding and Stage 3 full-text coding.

Operational mode

Automated candidate manifest plus human checklist completion.

Unit of analysis

One Stage 3 candidate record and its retrieval state.

Provenance basis

Manifest-generation logic in stage3_prep.py and the generated manual relevance checklist.

Research questions

Retrieval and adjudication gate that protects RQ2 and RQ3 corpus quality

Tagline

Automated candidate manifest plus human checklist completion

Canonical Source Paths

- `src/bps_review/extraction/stage3_prep.py`
- `src/review_stages/04_extraction/forms/stage3_manual_relevance_checklist.csv`
- `src/review_stages/04_extraction/outputs/stage3_candidate_manifest.csv`
- `src/review_stages/04_extraction/outputs/stage3_manual_fulltext_queue.csv`
- `src/review_stages/04_extraction/outputs/stage3_retrieval_validation.csv`

Purpose

This scheme governs the transition from Stage 2 abstract coding to Stage 3 full-text work. It standardizes which candidate reviews need manual retrieval, which records require manual relevance adjudication, and how retrieval status, risk signals, and reviewer decisions are recorded before deep coding begins.

It is not a logistics file. It is a standardized adjudication framework that decides whether a review can enter Stage 3, whether the retrieved text is adequate, and how problematic or ambiguous records are escalated for human judgment.

Manifest Fields

Automated fields written per candidate.

fulltext_status

Operational retrieval status of the candidate.

- **manual_retrieval_required** No machine-available full text was secured; human retrieval is needed.
- **pmc_open_available_not_cached** A PMCID exists but the file was not yet cached.
- **pmc_fulltext_cached** A PMC full text is already cached locally.
- **pmc_fulltext_fetched** A PMC full text was retrieved during the current run.
- **pmc_linked_fetch_failed** A PMCID link existed but retrieval or XML parsing failed.
- **pmc_fulltext_low_content_manual_check** Text was fetched but appears too short and needs manual inspection.

retrieval_source

Source of the PMCID or full-text link (existing metadata, PubMed elink, or Europe PMC).

fulltext_word_count

Word count of cached text, used to detect low-content full texts. (*free text*)

manual_retrieval_needed

Whether a human must retrieve the text.

- **yes** Retrieval required.
- **no** Machine text is available.

manual_relevance_priority

Adjudication urgency.

- **high** Withdrawal or missing pain focus.
- **medium** Other concern flags, or retrieval required with otherwise low signal.
- **low** No concern flags.

manual_relevance_flags

Pipe-delimited signal list describing why adjudication is needed. See the signal logic below. (*free text*)

osf_manual_adjudication_required

Currently set to yes for the checklist workflow.

cached_xml_path / cached_text_path

Relative cache paths for machine-fetched PMC files.

pubmed_url

Direct URL to the source record when a PMID is available.

Manual Relevance Signal Logic

Each signal is derived from title and abstract and drives priority.

withdrawn_or_retracted_signal

Assigned when the title or abstract indicates a withdrawn or retracted record. Forces high priority.

pain_focus_not_explicit

Assigned when pain is not explicit in title or abstract. Forces high priority.

chronicity_not_explicit

Assigned when chronic, persistent, or long-term framing is absent.

review_design_unclear

Assigned when review design is missing or too vague.

Reviewer Checklist Fields

Human adjudication fields completed against the manifest.

reviewer_decision

Human reviewer decision after checking the candidate. *(free text)*

reviewer_notes

Notes supporting the reviewer decision. *(free text)*

adjudication_decision

Final adjudicated status after disagreement resolution. *(free text)*

adjudication_notes

Adjudication rationale. *(free text)*

Proposed Refinement: Explicit Retrieval-Adequacy Gate

PROPOSED REFINEMENT

Awaiting expert evaluation. Not yet applied to the pipeline.

Refinement proposed for expert evaluation. The low-content flag currently relies on an implicit word-count heuristic. We propose a documented adequacy gate with a stated threshold (for example a minimum body word count and the presence of a results or discussion section) and an explicit adequacy value: adequate, borderline, or inadequate. Borderline and inadequate texts route to manual retrieval even when a file exists.

We also propose a reproducible priority matrix that maps the set of triggered signals to a priority level deterministically, so two coders reading the same manifest row derive the same priority. This removes the current tie-break ambiguity when several medium signals co-occur.

Test-Run Corpus State

Counts from the current test run, for context only.

Stage 2 included records

111

Stage 3 candidate manifest

92

Manual full-text queue

63

Reliability subset (capped at 20)

18

Primary Outputs Using This Scheme

- `src/review_stages/04_extraction/outputs/stage3_candidate_manifest.csv`
- `src/review_stages/04_extraction/forms/stage3_manual_relevance_checklist.csv`
- `src/review_stages/04_extraction/outputs/stage3_manual_fulltext_queue.csv`
- `src/review_stages/04_extraction/outputs/stage3_retrieval_validation.csv`